Footnote[1]

From “is this even true?” to “here’s what I can do about it.”

Live now: footnote.alessandrarye.com

Civic grounding for public claims: sources, transparency, and a real local next step.

Alee · Metana Full Stack Capstone · Final presentation, October 2026

The gap between seeing and doing

A resident sees a post: *“the new levy money is going to administrator raises.”*

Today she has three dead ends. She can’t easily check it. If she checks it, she gets a verdict with no reasoning. And nothing tells her what to actually do next, so she shares it, or shrugs.

Early signal: an attorney, a librarian, and a student, interviewed separately, each described this exact gap, and all three responded positively to Footnote.

1. **Verification tools stop at the verdict**
   Fact-checks and trust ratings end with “true-ish.” Then nothing.

2. **Civic software is sold to governments**
   The municipality is the customer; the resident is the audience.

3. **Community platforms have neither**
   Where claims actually spread, there’s no verification and no path to action.

Nobody connects *“I saw this claim”* to **“here is what I can do about it.”**

The solution: one path, four steps

1

Bring a claim

Paste the claim, pick a place. No login, no feed. Location is declared, never detected.

→

2

See the sources

A source card built from primary records, under a visible citation policy. Evidence, not verdict badges.

→

3

Open the record

The transparency layer: what was searched, used, rejected, and where AI was involved.

→

4

Get the next step

The upcoming meetings, the responsible officials, the organizations already on it, and where to find the records.

Then the session ends and evaporates. No profiles, no tracking, no engagement optimization, no ads. The constraints are the product.

It’s live, on real records

THE CLAIM YOU BROUGHT

"The levy money is going to classrooms."

What you're holding: a checkable statement of fact, about levy fund allocation.

WHAT THE RECORDS SAY

In fiscal year 2025, Olentangy Local School District spent $265,380,561 on instruction out of total general fund expenditures of $378,716,930. A ballot measure proposed levying an additional property tax estimated to collect $20,357,000 annually at a rate not exceeding 3 mills for each $1 of taxable value to pay current operating expenses. The Board of Education approved a Financial Forecast at a regular meeting on August 27, 2026. The records are presented as written; no verdict is offered.

Annual Comprehensive Financial Report FY2025, Olentangy Local School District primary

General fund expenditures for the fiscal year ended June 30, 2025: Instruction $265,380,561 of total expenditures of $378,716,930, with support services at $102,665,209.

Verified 2026-09-30

Live site, canonical claim, captured 2026-10-07

70% of general fund spending went to instruction in FY2025: $265.4M of $378.7M, audited.

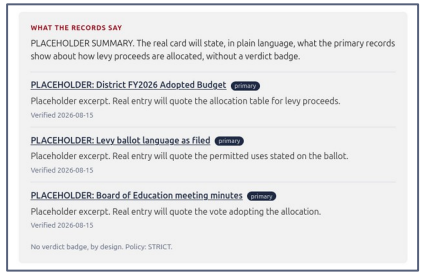
Three primary records

Audited financial report, FY2025

Ballot language, March 2024

Board meeting minutes, August 2026

Each links to the document itself and shows its verified date.



In August
this card was all placeholders. Now every fact is real, sourced, and dated.

How a claim moves through Footnote

“The levy money is going to classrooms”
+ Delaware County, Ohio

1 Classify AI

Claude Haiku labels the claim factual, opinion, or mixed. Strict JSON, validated in code.

2 Source

Picks the hand-verified record set that matches the claim, or none if nothing matches.

3 Policy AI

Keeps or rejects each source by tier. Haiku drafts the summary; rules check every sentence.

4 Invite

Meetings, officials, organizations, and resources, closest to the claim first.

Transparency record written by every stage: model and latency, sources kept and rejected, sentences kept and dropped, verified dates

Source card

What the records say, cited to primary sources.

Civic invitation

Who to contact and where to show up.

Transparency panel

How the answer was assembled, stage by stage.

Then the session evaporates. No profiles, no tracking, no ads.

The data: one county, by hand

18 officials

4 public meetings

4 organizations

6 resources

Every entry carries: its source link, the date it was verified, and its tier. Staleness is displayed, never hidden. The git history is the audit log.



Live site, 2026-10-07. The school board leads because the claim is about the levy.

The promise is not "the data is true." It's "you can always see where it came from and when it was checked."

The model writes, code decides

1 — The draft

Haiku gets the claim and the kept sources, and returns sentences, each tagged with the source it cites.

2 — The first run

It wrote that voters “approved” the levy, citing the ballot. The ballot lists the question, not the result.

3 — The fix: four rules, plain code

Cites a kept source · every figure appears in that source · no verdict language · no outcome word the source does not use.

Checker test run: six sentences, two kept, four dropped

KEPT	“The district spent $265,380,561 on instruction.”	*cites the financial report · figure found in the source*
DROPPED	“The district spent $300,000,000 on instruction.”	*cites the financial report · figure not in source*
DROPPED	“Voters approved the levy.”	*cites the ballot · outcome not stated in source*
KEPT	“The board approved the financial forecast.”	*cites the board minutes · the minutes use that word*
DROPPED	“This proves the levy money is misused.”	*contains verdict language*
DROPPED	“The county reported a surplus.”	*cites source 7 · does not point to a kept source*

The model writes the summary. Code decides what you see.

Tech stack: what is running

FE	**Frontend**	React 19 (Vite) · single-page web app for phone and desktop
API	**Backend**	Node.js + Express 5 REST API · four pipeline stages as separate modules (classify → source → policy → invite)
DATA	**Data**	Civic locality package as version-controlled YAML · every entry has a source, a tier, and a verified date · sessions are not stored
AI	**Model layer**	Claude Haiku 4.5 in two stages, classification and summary draft · one swappable adapter · clean fallback · every call logged to the transparency record
OPS	**Infra & process**	AWS EC2 + nginx + pm2, HTTPS via Let's Encrypt · pull requests into a protected main · dated decisions log

Changed since August: no PostgreSQL and no CI pipeline yet. Nothing about a session is kept, so there is nothing to store. Both are on the roadmap.

Your feedback, and honest limits

1. **A diagram of how a claim flows** — DONE
   In this deck and on the repo front page.
2. **Build out the data for one locality** — DONE
   Delaware County v1, plus three real source records.
3. **Keep using GitHub** — DONE
   Branch, pull request, protected main, decisions log.
4. **More interviews** — IN PROGRESS
   Round 1: attorney, librarian, student. Round 2 is underway.
5. **README, server, interface polish** — DONE
   Instructor feedback in October. All three are live.

What it cannot do yet

- The checker verifies attribution, figures, and wording. It does not yet verify that a paraphrase keeps the source's meaning.
- It has records for one question, the levy. Any other claim, even one its own data could answer, gets a "not yet covered" card until that context is built out.
- One county, maintained by hand.

A tool about trust has to be honest about its own edges.

Beyond the capstone: what comes next

1	**Next in the build**	A checker that catches the right number attached to the wrong idea · automated tests for every rule · sources connected to the claims people actually bring.
2	**Local stewards**	Residents and libraries in each place keep its information current, with the same source-and-date receipt on every entry. Who does this, and how, is the open question.
3	**Community source submission**	People can hand Footnote evidence; nobody can hand it an opinion. Reviewed pipeline, no posts, no identities, no audience.
4	**Topic packages**	The same provenance engine aimed at issues residents feel but don’t research: data-center siting, water and power agreements. Cite, never characterize.

Sustainability: free for citizens, always · institutional licensing (libraries, schools) · engine as an API · grants bridge · no ads, no engagement monetization, ever.

Footnote[1]

You pick a place. It shows you the records, shows its work, and hands you the next step. Then it forgets you.

Live: footnote.alessandrarye.com **Repo:** github.com/Alessandrarye/footnote

---

Questions welcome, especially the hard ones. The transparency record applies to me too.

Backup: competitors

Media trust tools

Ground News · NewsGuard · AI verification APIs (Factiverse, ClaimBuster)

Stop at the verdict.

Gov-side civic platforms

CivicPlus · OpenGov · Granicus

Sold to governments; the citizen isn't the customer.

Community platforms

Nextdoor · Facebook groups · local forums

Own the attention where claims spread. No verification, no path to action.

Google: substitutes for the verification half. Ran a Civic Information API for local officials data and shut it down in 2025.

OCLC: budget-adjacent in the library market, not functional. Footnote ingests only government primary records.

Footnote occupies the connection: verification → local action.

Backup: checker test output

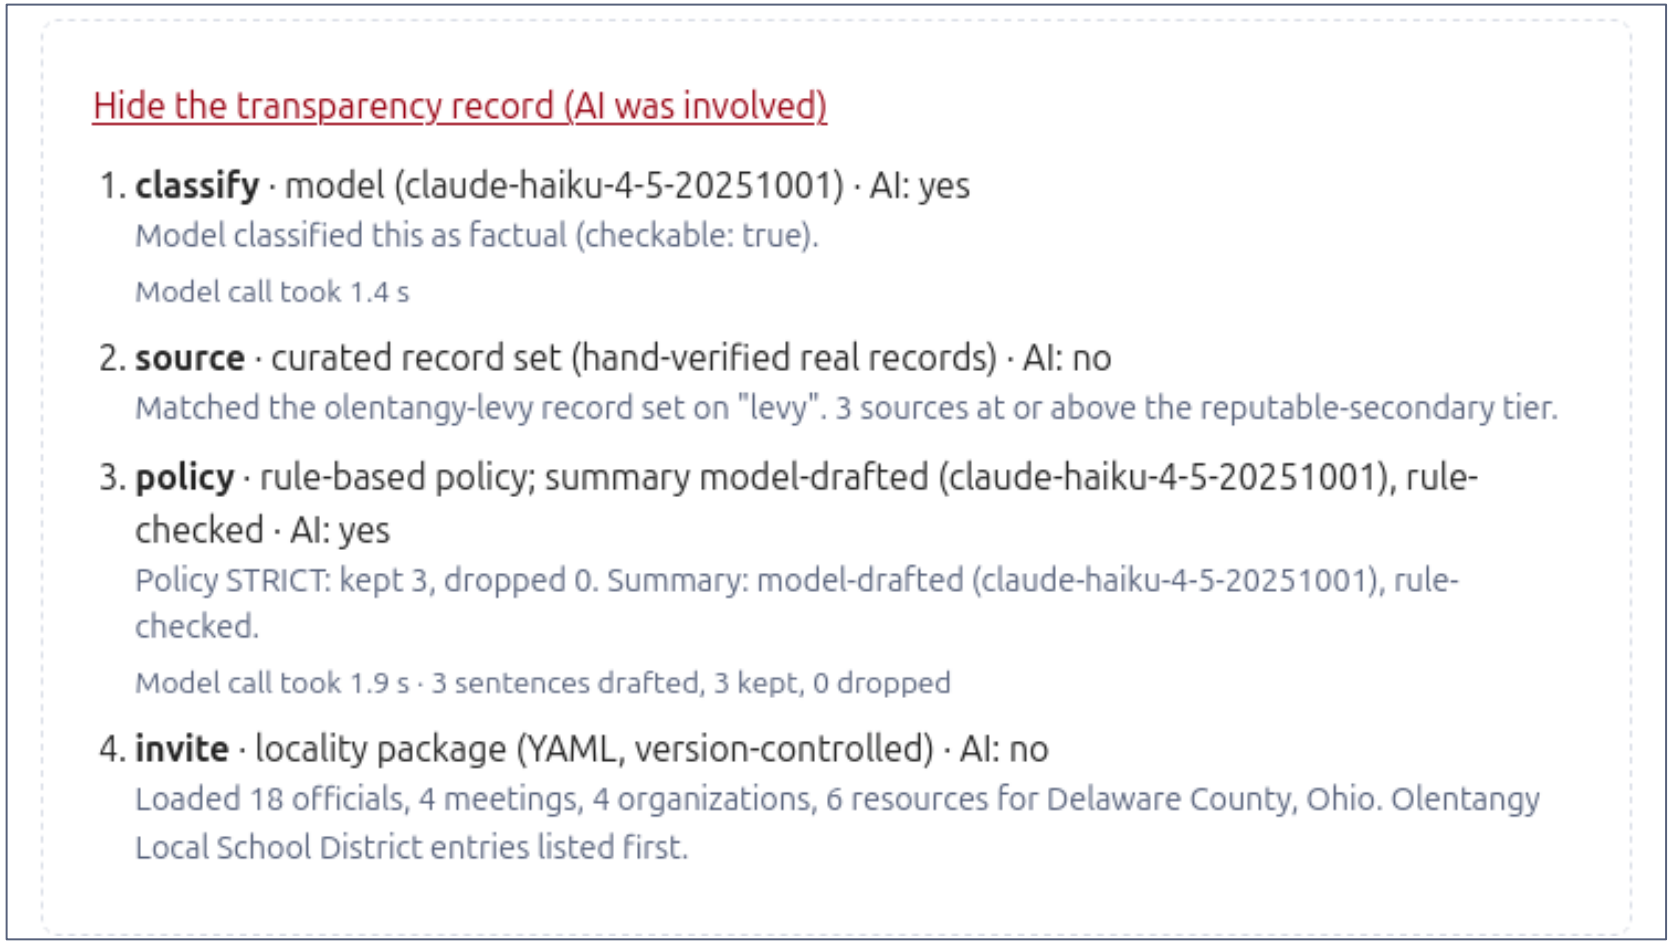
```
sandra@foundry:~/dev/footnote$ node <<'EOF'
const { checkSentence } = require("./server/pipeline/policy.js");
const { gatherSources } = require("./server/pipeline/source.js");
const kept = gatherSources("x", {}, { transparency: [] });

const tests = [
  ["real figure, source 1", { text: "The district spent $265,380,561 on instruction.", source: 1 }],
  ["made-up figure", { text: "The district spent $300,000,000 on instruction.", source: 1 }],
  ["outcome on the ballot", { text: "Voters approved the levy.", source: 2 }],
  ["outcome on the minutes", { text: "The board approved the financial forecast.", source: 3 }],
  ["verdict language", { text: "This proves the levy money is misused.", source: 1 }],
  ["source that does not exist", { text: "The county reported a surplus.", source: 7 }],
];

for (const [label, sentence] of tests) {
  const reason = checkSentence(sentence, kept);
  console.log(`${reason ? "DROPPED" : "KEPT   "}  ${label}${reason ? "  ->  " + reason : ""}`);
}
EOF
KEPT     real figure, source 1
DROPPED  made-up figure  ->  figure not in source: 300000000
DROPPED  outcome on the ballot  ->  outcome not stated in source: approv
KEPT     outcome on the minutes
DROPPED  verdict language  ->  contains verdict language
DROPPED  source that does not exist  ->  does not point to a kept source
sandra@foundry:~/dev/footnote$
```

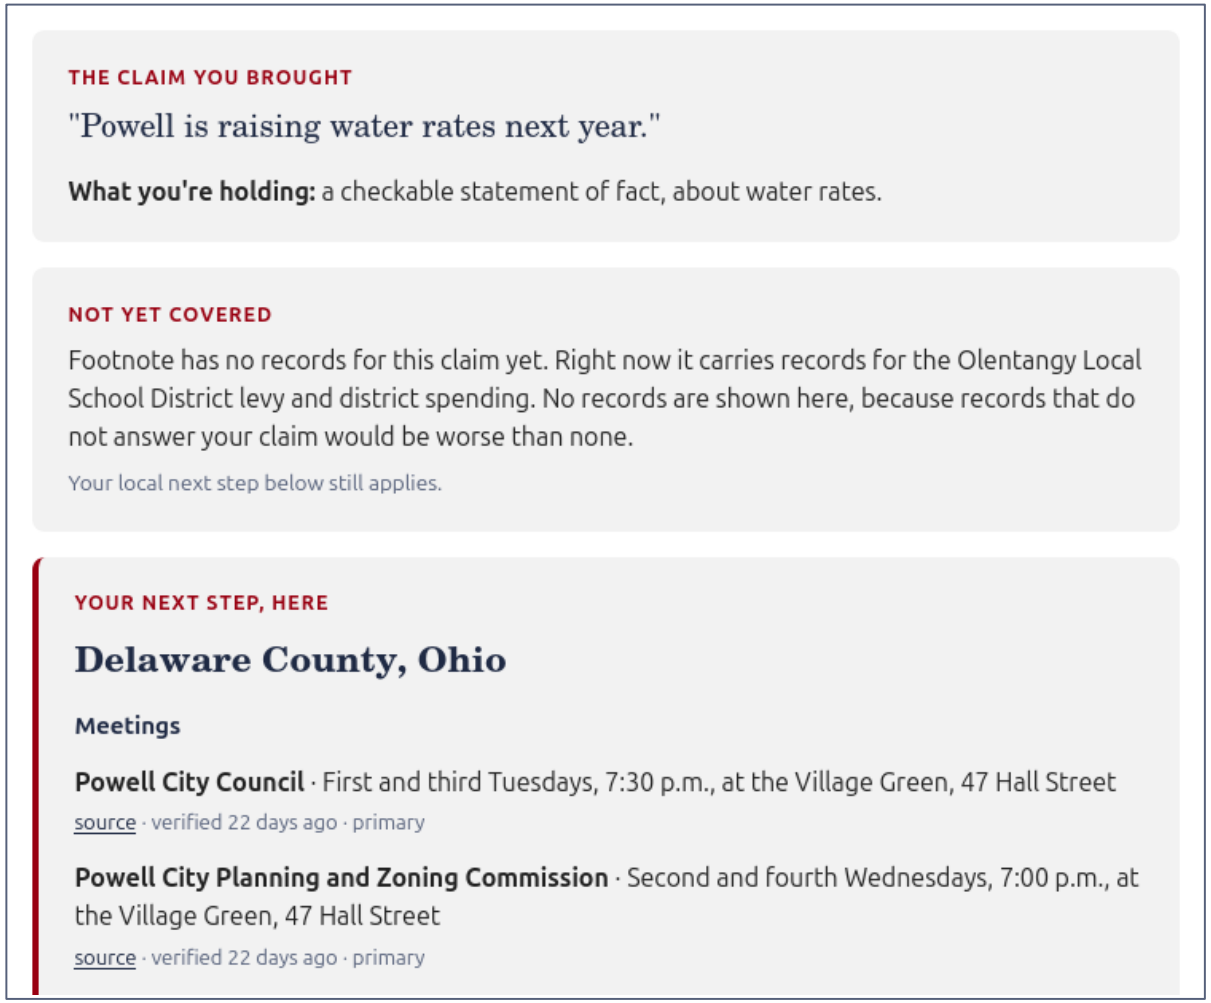
Run from ~/dev/footnote on 2026-10-04. The last six lines are the results.

Backup: the transparency record



Live site, 2026-10-07, canonical claim. Every stage says what it did, and whether AI was involved.

Backup: a claim it has no records for



Live site, 2026-10-07

No records beat the wrong records

Before this, every claim was shown the levy records. Now a claim outside the record set gets a plain statement of what Footnote does carry.

The next step still applies

The civic invitation still appears, in the locality's normal order, because the meetings and officials are real either way.

And it is on the record

The transparency record says no record set matched and that no summary was written. The model is not asked to fill the gap.